

## 导入数据

```
import torchvision
import torch
from torch.utils.data import DataLoader
import torch.nn.functional as F
import matplotlib.pyplot as plt
train_data = torchvision.datasets.MNIST(root="./data", train=True, download=False)
test_data = torchvision.datasets.MNIST(root="./data", train=False, download=False)

train_data[50][0]
```

```
/home/melancholy/Software/miniconda3/envs/pytorch/lib/python3.9/site-
packages/tqdm/auto.py:22: TqdmWarning: IPProgress not found. Please update jupyter
and ipywidgets. See
https://ipywidgets.readthedocs.io/en/stable/user_install.html
from .autonotebook import tqdm as notebook_tqdm
```

## 3

## img转换为tensor

```
# 数据集载入
train_data =
torchvision.datasets.MNIST(root="./data", train=True, transform=torchvision.transf
orms.ToTensor(), download=False)
test_data =
torchvision.datasets.MNIST(root="./data", train=False, transform=torchvision.trans
forms.ToTensor(), download=False)
train_loader = DataLoader(train_data, batch_size=32, shuffle=True)
test_loader = DataLoader(test_data, batch_size=32, shuffle=False)
```

## CNN网络层

```
class Net(torch.nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.con1 = torch.nn.Conv2d(1, 10, kernel_size=5) # 卷积层1
        self.con2 = torch.nn.Conv2d(10, 20, kernel_size=5) # 卷积层2
        self.pooling = torch.nn.MaxPool2d(2) # 池化层
        self.fc = torch.nn.Linear(320, 10) # 全链接层

# 前向传播
```

```

def forward(self,x):
    batch_size = x.size(0)
    x = F.relu(self.pooling(self.con1(x))) # 卷积1 + 池化
    x = F.relu(self.pooling(self.con2(x))) # 卷积2 + 池化
    x = x.view(batch_size,-1) # 铺平
    x = self.fc(x) # 全链接层
    return x

#模型实例化
model = Net()

device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
model.to(device)

```

```

Net(
  (con1): Conv2d(1, 10, kernel_size=(5, 5), stride=(1, 1))
  (con2): Conv2d(10, 20, kernel_size=(5, 5), stride=(1, 1))
  (pooling): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
    ceil_mode=False)
  (fc): Linear(in_features=320, out_features=10, bias=True)
)

```

## 定义损失函数交叉熵和优化其SGD

```

# lossfunction 损失函数使用交叉熵作为评判标准
lossfun = torch.nn.CrossEntropyLoss()

# 随机梯度下降
opt = torch.optim.SGD(model.parameters(),lr=0.01,momentum=0.5)

loss_all = []

```

## 训练函数，卷积，每隔200个数据记录一次loss

```

def train(epoch): # 训练模型
    running_loss = 0.0
    for i,(inputs,targets) in enumerate(train_loader,0):
        inputs,targets = inputs.to(device),targets.to(device) # to gpu
        opt.zero_grad() #清空 grad数据
        outputs = model(inputs) # 神经网络
        loss = lossfun(outputs,targets) # 求解loss
        loss.backward() # loss反向传播损失
        opt.step() # SGD---梯度下降
        running_loss += loss.item()
        if i % 200 == 0:
            loss_all.append(running_loss/200)
            print('[epoch:%d, %d] loss:%.3f' % (epoch+1,i+1,running_loss/200)) #
打印训练过程
            running_loss = 0.0

```

```

accuracy = []
def test():
    total = 0
    correct = 0
    with torch.no_grad(): # 关闭梯度下降
        for (inputs,targets) in test_loader:
            # to GPU
            inputs, targets = inputs.to(device), targets.to(device)
            # pred
            outputs = model(inputs)
            _,predicted = torch.max(outputs.data,dim=1)
            total += targets.size(0)
            correct += (predicted == targets).sum().item()
    accuracy.append(correct)
    print(100*correct/total)

```

```

if __name__ == '__main__':
    for epoch in range(20):
        train(epoch)
        test()

```

```

[epoch:1, 1] loss:0.001
[epoch:1, 201] loss:0.024
[epoch:1, 401] loss:0.026
[epoch:1, 601] loss:0.026
[epoch:1, 801] loss:0.023
[epoch:1, 1001] loss:0.021
[epoch:1, 1201] loss:0.023
[epoch:1, 1401] loss:0.025
[epoch:1, 1601] loss:0.024
[epoch:1, 1801] loss:0.024
98.89
[epoch:2, 1] loss:0.000
[epoch:2, 201] loss:0.022
[epoch:2, 401] loss:0.021
[epoch:2, 601] loss:0.024
[epoch:2, 801] loss:0.022
[epoch:2, 1001] loss:0.017
[epoch:2, 1201] loss:0.025
[epoch:2, 1401] loss:0.027
[epoch:2, 1601] loss:0.022
[epoch:2, 1801] loss:0.025
98.95
[epoch:3, 1] loss:0.000
[epoch:3, 201] loss:0.021
[epoch:3, 401] loss:0.028
[epoch:3, 601] loss:0.023
[epoch:3, 801] loss:0.022
[epoch:3, 1001] loss:0.020

```

```
[epoch:3, 1201] loss:0.022
[epoch:3, 1401] loss:0.023
[epoch:3, 1601] loss:0.025
[epoch:3, 1801] loss:0.021
98.83
[epoch:4, 1] loss:0.000
[epoch:4, 201] loss:0.019
[epoch:4, 401] loss:0.022
[epoch:4, 601] loss:0.024
[epoch:4, 801] loss:0.025
[epoch:4, 1001] loss:0.022
[epoch:4, 1201] loss:0.016
[epoch:4, 1401] loss:0.024
[epoch:4, 1601] loss:0.021
[epoch:4, 1801] loss:0.024
98.83
[epoch:5, 1] loss:0.000
[epoch:5, 201] loss:0.018
[epoch:5, 401] loss:0.017
[epoch:5, 601] loss:0.022
[epoch:5, 801] loss:0.024
[epoch:5, 1001] loss:0.021
[epoch:5, 1201] loss:0.016
[epoch:5, 1401] loss:0.022
[epoch:5, 1601] loss:0.024
[epoch:5, 1801] loss:0.023
98.93
[epoch:6, 1] loss:0.000
[epoch:6, 201] loss:0.015
[epoch:6, 401] loss:0.019
[epoch:6, 601] loss:0.019
[epoch:6, 801] loss:0.021
[epoch:6, 1001] loss:0.023
[epoch:6, 1201] loss:0.022
[epoch:6, 1401] loss:0.023
[epoch:6, 1601] loss:0.020
[epoch:6, 1801] loss:0.017
98.72
[epoch:7, 1] loss:0.000
[epoch:7, 201] loss:0.014
[epoch:7, 401] loss:0.018
[epoch:7, 601] loss:0.019
[epoch:7, 801] loss:0.020
[epoch:7, 1001] loss:0.020
[epoch:7, 1201] loss:0.024
[epoch:7, 1401] loss:0.023
[epoch:7, 1601] loss:0.020
[epoch:7, 1801] loss:0.019
98.89
[epoch:8, 1] loss:0.000
[epoch:8, 201] loss:0.015
[epoch:8, 401] loss:0.018
[epoch:8, 601] loss:0.021
[epoch:8, 801] loss:0.017
[epoch:8, 1001] loss:0.018
```


```
[epoch:8, 1201] loss:0.018
[epoch:8, 1401] loss:0.024
[epoch:8, 1601] loss:0.018
[epoch:8, 1801] loss:0.021
98.74
[epoch:9, 1] loss:0.000
[epoch:9, 201] loss:0.019
[epoch:9, 401] loss:0.024
[epoch:9, 601] loss:0.018
[epoch:9, 801] loss:0.019
[epoch:9, 1001] loss:0.017
[epoch:9, 1201] loss:0.018
[epoch:9, 1401] loss:0.013
[epoch:9, 1601] loss:0.020
[epoch:9, 1801] loss:0.014
98.84
[epoch:10, 1] loss:0.000
[epoch:10, 201] loss:0.016
[epoch:10, 401] loss:0.018
[epoch:10, 601] loss:0.015
[epoch:10, 801] loss:0.019
[epoch:10, 1001] loss:0.017
[epoch:10, 1201] loss:0.018
[epoch:10, 1401] loss:0.022
[epoch:10, 1601] loss:0.013
[epoch:10, 1801] loss:0.016
98.8
[epoch:11, 1] loss:0.000
[epoch:11, 201] loss:0.018
[epoch:11, 401] loss:0.022
[epoch:11, 601] loss:0.019
[epoch:11, 801] loss:0.018
[epoch:11, 1001] loss:0.014
[epoch:11, 1201] loss:0.018
[epoch:11, 1401] loss:0.018
[epoch:11, 1601] loss:0.012
[epoch:11, 1801] loss:0.013
99.01
[epoch:12, 1] loss:0.000
[epoch:12, 201] loss:0.013
[epoch:12, 401] loss:0.014
[epoch:12, 601] loss:0.018
[epoch:12, 801] loss:0.017
[epoch:12, 1001] loss:0.016
[epoch:12, 1201] loss:0.012
[epoch:12, 1401] loss:0.015
[epoch:12, 1601] loss:0.019
[epoch:12, 1801] loss:0.020
98.84
[epoch:13, 1] loss:0.000
[epoch:13, 201] loss:0.013
[epoch:13, 401] loss:0.015
[epoch:13, 601] loss:0.020
[epoch:13, 801] loss:0.013
[epoch:13, 1001] loss:0.016
```

[epoch:13, 1201] loss:0.018  
[epoch:13, 1401] loss:0.012  
[epoch:13, 1601] loss:0.018  
[epoch:13, 1801] loss:0.014  
98.59  
[epoch:14, 1] loss:0.000  
[epoch:14, 201] loss:0.014  
[epoch:14, 401] loss:0.015  
[epoch:14, 601] loss:0.016  
[epoch:14, 801] loss:0.015  
[epoch:14, 1001] loss:0.015  
[epoch:14, 1201] loss:0.017  
[epoch:14, 1401] loss:0.012  
[epoch:14, 1601] loss:0.015  
[epoch:14, 1801] loss:0.013  
98.89  
[epoch:15, 1] loss:0.000  
[epoch:15, 201] loss:0.014  
[epoch:15, 401] loss:0.017  
[epoch:15, 601] loss:0.012  
[epoch:15, 801] loss:0.013  
[epoch:15, 1001] loss:0.018  
[epoch:15, 1201] loss:0.018  
[epoch:15, 1401] loss:0.016  
[epoch:15, 1601] loss:0.017  
[epoch:15, 1801] loss:0.013  
99.0  
[epoch:16, 1] loss:0.000  
[epoch:16, 201] loss:0.009  
[epoch:16, 401] loss:0.013  
[epoch:16, 601] loss:0.015  
[epoch:16, 801] loss:0.018  
[epoch:16, 1001] loss:0.013  
[epoch:16, 1201] loss:0.014  
[epoch:16, 1401] loss:0.011  
[epoch:16, 1601] loss:0.012  
[epoch:16, 1801] loss:0.020  
98.93  
[epoch:17, 1] loss:0.000  
[epoch:17, 201] loss:0.016  
[epoch:17, 401] loss:0.012  
[epoch:17, 601] loss:0.018  
[epoch:17, 801] loss:0.015  
[epoch:17, 1001] loss:0.012  
[epoch:17, 1201] loss:0.013  
[epoch:17, 1401] loss:0.012  
[epoch:17, 1601] loss:0.013  
[epoch:17, 1801] loss:0.013  
98.89  
[epoch:18, 1] loss:0.000  
[epoch:18, 201] loss:0.015  
[epoch:18, 401] loss:0.011  
[epoch:18, 601] loss:0.009  
[epoch:18, 801] loss:0.015  
[epoch:18, 1001] loss:0.011

```
[epoch:18, 1201] loss:0.013
[epoch:18, 1401] loss:0.016
[epoch:18, 1601] loss:0.014
[epoch:18, 1801] loss:0.013
98.99
[epoch:19, 1] loss:0.000
[epoch:19, 201] loss:0.010
[epoch:19, 401] loss:0.014
[epoch:19, 601] loss:0.014
[epoch:19, 801] loss:0.013
[epoch:19, 1001] loss:0.013
[epoch:19, 1201] loss:0.011
[epoch:19, 1401] loss:0.011
[epoch:19, 1601] loss:0.012
[epoch:19, 1801] loss:0.014
98.89
[epoch:20, 1] loss:0.000
[epoch:20, 201] loss:0.015
[epoch:20, 401] loss:0.009
[epoch:20, 601] loss:0.012
[epoch:20, 801] loss:0.012
[epoch:20, 1001] loss:0.014
[epoch:20, 1201] loss:0.013
[epoch:20, 1401] loss:0.010
[epoch:20, 1601] loss:0.015
[epoch:20, 1801] loss:0.009
99.02
```

## 打印loss


```
# plot loss
x = [i for i in range(len(loss_all))]
fig = plt.figure(dpi=300)
plt.plot(x, loss_all)
plt.xlabel("iter")
plt.ylabel("loss")
plt.show()
```



## 打印test accuracy

```
# plot accuracy
y = [i for i in range(len(accuracy))]
fig = plt.figure(figsize=(10,5),dpi=300)
for i in accuracy:
    i/=100
plt.plot(y,accuracy)
plt.xlabel("epoch")
plt.ylabel("accuracy")
plt.title("accuracy-epoch")
plt.show()
```





## 预测及可视化参照

```
count = 10
with torch.no_grad(): # 关闭梯度下降
    for (inputs,targets) in test_loader:
        # to GPU
        inputs, targets = inputs.to(device), targets.to(device)
        # pred
        outputs = model(inputs)
        _,predicted = torch.max(outputs.data,dim=1)
        for i in range(6):
            plt.subplot(2,3,i+1)
            plt.tight_layout()
            plt.imshow(test_data[i][0],cmap="gray",interpolation='none')
            plt.title("Prediction: {}".format(predicted[i]))
            plt.xticks([])
            plt.yticks([])
        break
```

Prediction: 7



Prediction: 2



Prediction: 1



Prediction: 0



Prediction: 4



Prediction: 1

